

QUANTIZED NEURAL NETWORK WITH FOUR BITS COMPRESSION

Ladislav de Naurois, Matteo Turchetta, Luca Corinzia

Department of Computer Science
ETH Zürich
Zürich, Switzerland

ABSTRACT

Neural Networks are a class of models for making inference over complex non-linear functions that have established the state of the art for several machine learning tasks. Despite their success, their diffusion in the world of embedded devices is limited by their memory and computational requirements. These requirements stem from the high number of parameters, usually stored as float values, that is required to represent a Neural Network. To overcome this, quantization compresses these parameters to a representation that uses a user-defined number of bits, k .

1. INTRODUCTION

In recent years we are witnessing an exponential increase in the amount of data available for analysis in almost all scientific disciplines. As a result, there is a high interest in machine learning methods to conduct such analysis. Among these, Neural Networks (NNs) are regarded as one of the most promising techniques. They have been successfully applied to a wide range of tasks including medical applications [1], image recognition [2] and robotics [3].

One of the main drawbacks of NNs is their high number of parameters. As a consequence, NNs require a lot of memory resources for storage and a lot of computational resources for training and forward prediction. While the training phase is usually performed on parallel computing architectures where there is an abundance of both computational and memory resources, the platforms where trained networks are deployed usually have more limited capabilities (e.g. mobile phones). This problem has steered attention of the research community toward reducing the memory and computation requirements for trained NNs.

A promising research direction is the one of quantized neural networks (QNNs). The central idea to QNNs is to compress the parameters of the network from their float representation to a light-weight one based on quantization bins. The parameter space is divided into a predefined number of bins and each parameter float value is mapped to a bin. The number of bins trades-off the accuracy versus the gain in memory and computation requirements.

In this work we present an optimized implementation of a QNN for the forward prediction on the MNIST data set that makes use of four bits quantization.

Related work. The research regarding NNs compression focuses mostly on memory requirements. Thus the performance benefits that can be obtained as a by-product of compression are almost unexplored. Hence contributions in the literature are mostly on the algorithmic side rather than on the code optimization one. For example, [4] use k -means clustering to reduce the size of a convolutional NN. The work of [5] exploits the over-parametrization of NNs by randomly grouping parameters by means of a hashing function. In [6] the authors reduce the size of the convolution filters of a NN using low-ranking approximation methods. More closely related to our method are the works that explicitly reduce the number of bits used to represent the weights in a NN. Among these [7] propose a high-accuracy four bits quantization scheme for recurrent NNs, a type of NN that is notorious for low prediction performance when quantized. However, their open source implementation is not optimized for performance. In [8] present a high-performance implementation of one bit QNNs optimized for Graphical Processing Units (GPUs) is presented. The authors of [9] introduce an implementation of an eight bits QNN for speech recognition optimized for CPUs.

Our contribution consists in an high-performance implementation of four bits QNN optimized for CPUs. While this level of quantization can yield substantial improvements in memory and computation requirements, it presents implementation challenges due to the byte addressability of most computer memories and to the lack of built-in data type for four bits integers.

2. BACKGROUND: NNS AND QNNS

In this section we formally introduce NNs and QNNs and relative notation.

Artificial neural networks. An artificial neural network (NN) is formally a (in general) non linear map from an input vector $\mathbf{x} \in \mathbb{R}^{d_i}$ to an output vector $f(\mathbf{x}) = \mathbf{y} \in \mathbb{R}^{d_o}$. The map f is built recursively applying at step t a linear

transformation $\mathbf{a}_t = \mathbf{W}_t \mathbf{x}_t + \mathbf{b}_t$ and a non-linear transformation $\mathbf{x}_{t+1} = \phi_t(\mathbf{a}_t)$. The matrix \mathbf{W}_t is called *weight matrix*, and the vector \mathbf{b}_t is called *bias vector*. The step t is also known as the *layer index*.

Quantized neural network. A quantized neural network (QNN) is a NN that uses low precision weight matrix and bias vector. Formally, given a NN with parameters $\{\mathbf{W}_t\}, \{\mathbf{b}_t\}$ and activation functions ϕ_t , the quantized implementation is the NN that apply at each layer the linear transformation $\mathbf{a}_t = \mathcal{Q}(\mathbf{W}_t)\mathcal{Q}(\mathbf{x}_t) + \mathcal{Q}(\mathbf{b}_t)$ and a non-linear transformation $\mathbf{x}_{t+1} = \phi_t(\mathbf{a}_t)$, where the function $\mathcal{Q}(\cdot)$ is introduced in the following.

Matrix quantization. For a matrix \mathbf{A} , the function $\mathcal{Q}(\mathbf{A})$ returns a low precision encoding of the matrix \mathbf{A} . It first computes the minimum (mn) and the maximum (mx) entry of the matrix \mathbf{A} , then given k bits it builds a linear binning of the continuous interval $[mn, mx]$ into 2^k many bins. The bin size of the quantization $\Delta(\mathbf{A})$ is then $\Delta(\mathbf{A}) = \frac{mx-mn}{2^k}$. To insure that the value 0 is represented exactly as a bin value, its index is computed as $z(\mathbf{A}) = \text{sat}([-mn/\Delta(\mathbf{A})])$, where the brackets $[\cdot]$ stand for the rounding to the closest integer and the $\text{sat}(\cdot)$ function saturates an integer value into the integer value representable with k bits, hence $\text{sat}(n) = \max(0, \min(n, 2^k - 1))$. The bin values are then $\{(i - z(\mathbf{A}))\Delta(\mathbf{A}), 0, \dots, 2^k - 1\}$. Then every entry A_{ij} is quantized to the closest bin value. The quantize matrix $\mathcal{Q}(\mathbf{A})$ and the quantized integer matrix $\tilde{\mathcal{Q}}(\mathbf{A})$ have respectively the bin value and the bin index as entry ij . Note that the matrix $\mathcal{Q}(\mathbf{A})$ is a real-valued matrix, while $\tilde{\mathcal{Q}}(\mathbf{A})$ is k -bit integer valued, and also that the following holds:

$$\mathcal{Q}(\mathbf{A}) = (\tilde{\mathcal{Q}}(\mathbf{A}) - z(\mathbf{A})\mathbf{J})\Delta(\mathbf{A}) \quad (1)$$

where \mathbf{J} is a matrix with all entries equal to one. The algorithm is showed in algorithm 1.

Algorithm 1 Quantize

- 1: compute $mn = \min A_{ij}$ and $mx = \max A_{ij}$
 - 2: $\Delta = \frac{mx-mn}{2^k}$.
 - 3: $z = -mn/\Delta$
 - 4: **for** $i, j = 1, \dots, N$ **do**
 - 5: $\tilde{\mathcal{Q}}(\mathbf{A})_{ij} = \text{saturate}([A_{ij}/\Delta + z])$
-

Quantized Matrix-Matrix Multiplication. Given two matrices \mathbf{L} and \mathbf{R} , we want to compute the product $\mathcal{Q}(\mathbf{L})\mathcal{Q}(\mathbf{R})$. Using eq. (1) and we write the product as

$$\mathcal{Q}(\mathbf{L})\mathcal{Q}(\mathbf{R}) = \Delta(\mathbf{L}) \left(\tilde{\mathcal{Q}}(\mathbf{L}) - z(\mathbf{L})\mathbf{J} \right) \left(\tilde{\mathcal{Q}}(\mathbf{R}) - z(\mathbf{R})\mathbf{J} \right) \Delta(\mathbf{R}) \quad (2)$$

so that we only need to perform MMM on k -bit integer valued matrix. Inverting the equation 1 we can then obtain the

k -bit integer valued product matrix as

$$\tilde{\mathcal{Q}}(\mathbf{LR}) = \text{sat}\left(\left[\frac{1}{\Delta(\mathbf{LR})} \mathcal{Q}(\mathbf{L})\mathcal{Q}(\mathbf{R}) + z(\mathbf{LR})\mathbf{J}\right]\right) \quad (3)$$

The algorithm is showed in algorithm 2.

Algorithm 2 QMMM

- 1: compute $\mathcal{Q}(\mathbf{L})\mathcal{Q}(\mathbf{R})$ as in eq. (2)
 - 2: compute the k -bit integer matrix $\tilde{\mathcal{Q}}(\mathbf{LR})$ as in eq. (3)
-

3. PERFORMED OPTIMIZATION

In this section we propose an optimized implementation of the quantization and Quantized Matrix-Matrix Multiplication (QMMM) functions introduced in section 2 that uses four bits compression. We start by presenting the data structure that we use for our baseline implementation. We continue by analysing the bottlenecks of this implementation and by proposing a set of solutions to achieve a higher performance.

Baseline implementation. We presented the algorithm at the base of a straight forward implementation of a simple QNN using a k -bits compression scheme in section 2. Here we introduce the data structure necessary to a naive implementation in case $k = 4$. Using a 4-bits compression scheme presents challenges due to the byte addressability of the computer memory and to the lack of a built-in 4-bits integer data type. As a consequence, we have to define our custom data structure. One way of operating on entities that require less than a byte for storage is to use structs in combination with bit fields. Nevertheless, byte addressability of the memory does not make it possible to load or store less than one byte at a time. Hence, using bit fields to define a custom data type that stores a single 4-bits integer is wasteful. In the proposed solution we define the `uint4x4_t` data structure. This is a 2 bytes struct that exploits bit fields to store four integers of 4 bits each. The reason why we pack four integers in one struct instead of two is because it allows us to define a data type that is more convenient for QMMM. Given a $n \times m$ weight matrix \mathbf{A} of floats, its 4-bits quantized counterpart $\tilde{\mathcal{Q}}(\mathbf{A})$ is stored as an $\frac{n}{2} \times \frac{m}{2}$ matrix of `uint4x4_t`. Because of this design choice, we assume both n and m to be even. In particular, the element (i, j) of $\tilde{\mathcal{Q}}(\mathbf{A})$ contains the quantized representation of the elements of \mathbf{A} at the following indices: $(2i, 2j), (2i, 2j + 1), (2i + 1, 2j), (2i + 1, 2j + 1)$. The relation between the logical layout and the memory layout for the original matrix \mathbf{A} and its quantized version that uses `uint4x4_t` data structure, $\tilde{\mathcal{Q}}(\mathbf{A})$, can be seen in fig. 1. The difference between the memory layout of the \mathbf{A} and $\tilde{\mathcal{Q}}(\mathbf{A})$ will play an important role in the vectorization of the code.

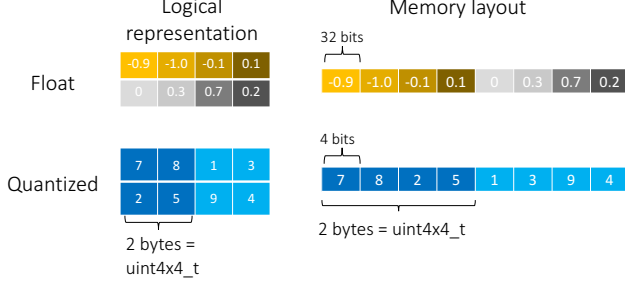


Fig. 1: Logical and memory layouts of a weight matrix and its corresponding 4-bits quantized version stored as a matrix of *uint4x4_t*. Cells that share the same color are stored in the same data structure.

Operation count optimization. The first optimization we present concerns the reduction of redundant computation. Using simple linear algebra, we can rewrite eq. (2) as:

$$\begin{aligned} \mathcal{Q}(\mathbf{L})\mathcal{Q}(\mathbf{R}) = & \Delta(\mathbf{L})\Delta(\mathbf{R})(\tilde{\mathcal{Q}}(\mathbf{L})\tilde{\mathcal{Q}}(\mathbf{R}) - \mathbf{z}(\mathbf{L})\mathbf{J}\tilde{\mathcal{Q}}(\mathbf{R}) + \\ & - \mathbf{z}(\mathbf{R})\tilde{\mathcal{Q}}(\mathbf{L})\mathbf{J} + \mathbf{z}(\mathbf{L})\mathbf{z}(\mathbf{R})\mathbf{J}\mathbf{J}). \end{aligned} \quad (4)$$

This formulation of the QMMM makes it easy to see its redundant computation. The result of the product $\mathbf{J}\tilde{\mathcal{Q}}(\mathbf{R})$ contained in the second term inside the parenthesis is a matrix that has all rows equal to each other. In particular the i^{th} term of any such row is the sum of the elements of the i^{th} column of $\tilde{\mathcal{Q}}(\mathbf{R})$. As a consequence, one can compute such row only once and save computation. In particular, this means we move ops from the inner most loop of the QMMM to the second inner most loop. Thus, the count of such ops grows quadratically with the size of the weight matrix rather than cubically. A similar reasoning can be applied to the third and fourth term inside the parenthesis.

Memory optimization. Another important optimization of our implementation is related to the efficient use of the memory. To increase the temporal locality of the QMMM kernel, we exploit the well known blocking strategy, both for cache and for register. The parameter of the blocking are the cache-block size N_b and the register-block size n_b , as outlined in fig. 2

Vectorization. Finally, a critical aspect of our implementation concerns the vectorization of the code. The gain in performance that can be expected vary depending on what type of data a function operates on. For example, if we consider the functions needed to implement the quantization step, we can expect to observe a gain in performance up to $\times 8$ because an AVX register can operate on 8 floats simultaneously. On the other hand, when we consider the QMMM, in principle we can expect a gain in performance up to $\times 16$. This is because, in order to have the representation power that is necessary to compute correctly the




Fig. 2: Blocking parameters for the QMMM

dot products that appear in the term $\tilde{\mathcal{Q}}(\mathbf{L})\tilde{\mathcal{Q}}(\mathbf{R})$ of eq. (4), we use 16-bits integers as accumulators. As a consequence, our implementation of quantization and QMMM is divided into sub-functions that try to keep as separate as possible the integer computation from the float computation. Among them, the most important are:

- *round-sat*: computes the round and saturation operation used in line 5 of algorithm 1 and line 2 of algorithm 2.
- *quantize*: given a weight matrix \mathbf{A} performs the operation necessary to compute $\tilde{\mathcal{Q}}(\mathbf{A})$ except for rounding and saturating (e.g. computing mn , mx , $\Delta(\mathbf{A})$).
- *qmmm_kernel*: computes $\tilde{\mathcal{Q}}(\mathbf{L})\tilde{\mathcal{Q}}(\mathbf{R})$, or, in other words, the dot products of the inner most loop of the QMMM.
- *trick*: computes the sum over columns and rows of $\tilde{\mathcal{Q}}(\mathbf{R})$ and $\tilde{\mathcal{Q}}(\mathbf{L})$ respectively. As explained before, this is needed to reduce the overall ops count of QMMM.
- *add_trick*: adds the appropriate element of the row and the columns computed by the *trick* function to the result of the dot product computed by *qmmm_kernel*.

Table 1 summarizes the cost in FLOPS or IOPS of each of the functions above.

Among the numerous advantages in terms of performance that come from using 4-bits compression, there are also drawbacks due to the memory layout of the *uint4x4_t* data structure shown in fig. 1. For example, listing 1 shows the overhead in terms of masking, shifting and blending that is required to load two rows of the quantized weight matrix.

Function	Ops count (int or float)
<i>round-sat</i>	$5N^2$ FLOPS
<i>quantize</i>	$7N^2$ FLOPS
<i>qmmm_kernel</i>	$2N^3$ IOPS
<i>trick</i>	$2N^2 + 2N$ IOPS
<i>add_trick</i>	$3N^3$ IOPS

Table 1: Cost analysis

```

1  __m256i tmp = _mm256_loadu_si256((__m256i const *)a);
2
3  // Mask for odd indices in memory
4  __m256i odd = _mm256_and_si256(tmp, _mm256_set1_epi8
5  (15));
6
7  // Mask for even indices in memory
8  __m256i even = _mm256_and_si256(tmp, _mm256_set1_epi8
9  (240));
10
11 // Shift and blend to recover rows
12 b_mask = _mm256_set1_epi16(32768);
13 *b1 = _mm256_blendv_epi8(odd, _mm256_srli_epi64(even,
14  4), b_mask);
15 *b2 = _mm256_blendv_epi8(_mm256_srli_epi64(odd, 8),
16  _mm256_srli_epi64(even, 4), b_mask);

```

Listing 1: Load of two *uint4x4_t* rows.

4. EXPERIMENTAL RESULTS

In this section we perform the empirical evaluation of the optimization performed as outlined in the section 3. Every code version is tested for correctness on a small size example with hand-computed output and on several large-size random instances with output given by the naive implementation.

Experimental setup. For the empirical evaluation of the code, we use an Skylake processor (3.5 GHz, L1 cache 128 KB, L2 cache 1 MB, L3 cache 6 Mb). The compiler used is g++ with flags "-O3 -fno-tree-vectorize -march=native -mavx". The matrix size varies in the range [30, 1000].

Results: operational count optimization. We first evaluate the gain in the runtime regarding the optimization in the operational count performed with the trick as in eq. (4). From fig. 3 we can see the decrease in performance of the trick version with respect to the naive implementation. fig. 4 shows that the operation count optimization gives an overall speed-up of 15%. In fig. 5 we can see the contribution of each function of the pipeline in the overall runtime. In the following the trick version is further optimized.

Results: blocking for MMM. The blocking parameter used is $N_b = 30$ for the cache blocking parameter and $n_b = 3$ for the register blocking parameter, as described in section 3. fig. 3 shows the gain in performance with respect to the naive.trick implementation. The large n performance gain is approximately 2X. Note that the blocking parameter for cache and for register are not fine tuned, so a further speedup could be possible.

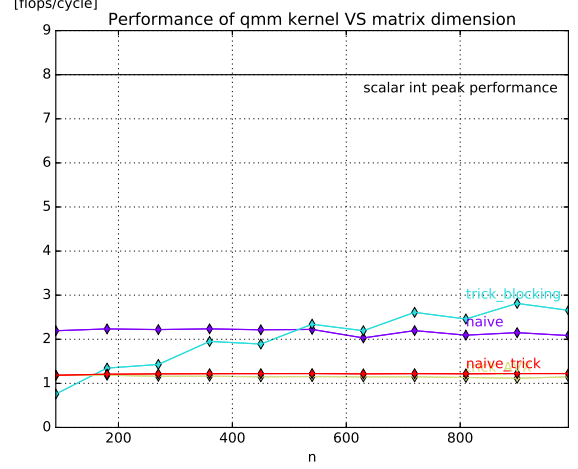


Fig. 3: Performance plot for the QMM kernel

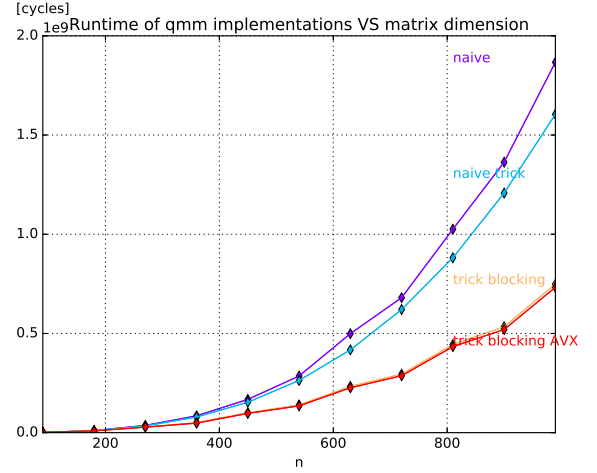


Fig. 4: Runtime plot of the overall pipeline

Results: vectorization. In this paragraph we evaluate the performance gain of the vectorization of the sub-functions *trick_vector*, *quantize*, *add_trick_vector* and *round_saturation* as outlined in section 3. fig. 6 and fig. 8 show a linear increase of the performance for small instances. This is due to a border effect. The instance size in the performance plot are not in general a multiple of the vector size (16X16 Bits), hence for small instance size the scalar computation can be non-negligible. Moreover the scalar contribution decreases linearly with the size of the instance.

The maximal gain in performance that the vectorization can allow for the integer functions *trick_vector* and *add_trick_vector* is 16X, that is the size of the accumulator vector. The measured performance gain are respectively 9.8X and 8.5X.

fig. 7 and fig. 9 show the performance plot for the functions *quantize* and *round_saturation*. The measured performance gain are respectively 9.2X and 14.2X.

In fig. 4 we can see the runtime plot for the whole pipeline.




Fig. 5: Contribution of each naive-version sub-function in the overall runtime.

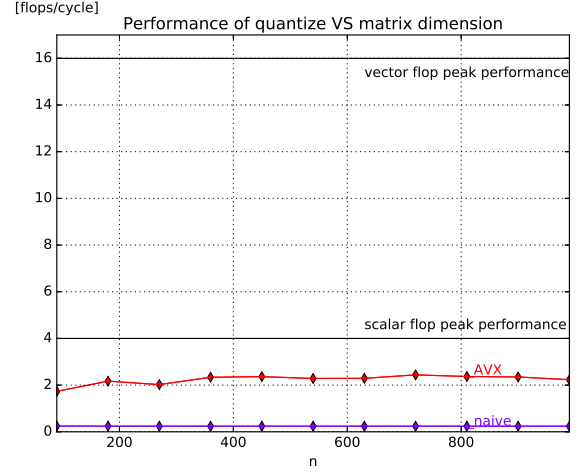


Fig. 7: Performance plot of the function *quantize*

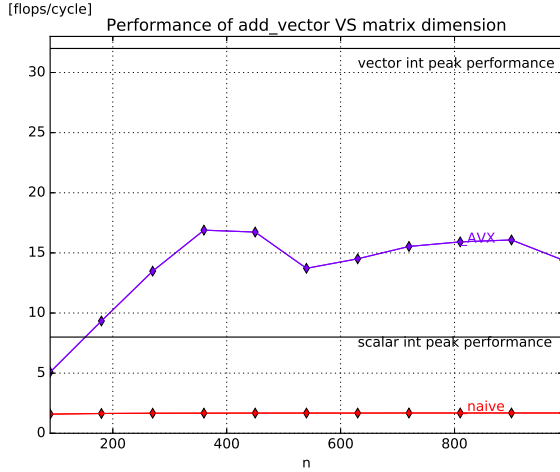


Fig. 6: Performance plot of the function *add_vector*




Fig. 8: Performance plot of the function *trick_vector*

The implementation *trick_blocking_AVX* is obtained combining the *QMM_kernel_blocking* with the vectorized implementation of all the other sub-functions. The overall speedup is 2.5X.

5. CONCLUSIONS

The experimental results presented in the previous sections outline the benefits of performing vectorization and matrix blocking to achieve a significant speedup over a naive implementation. The key ingredient in the approach outlined above lies in the decision to map a tile of 2×2 floats to a single 16-bit integer (using a struct of 4 members with 4 reserved bits respectively). Through this construction, the algorithms outlined in the previous are accessible in a convenient form that lends itself to simply vectorize the




Fig. 9: Performance plot of the function *round_saturation*

load, quantize, round and saturate functions, while also enabling a simplified blocking scheme for quantized matrix multiplication. More importantly, each of the individual optimizations performed demonstrate the ability to improve over the current optimized implementations of the General Matrix Multiplication with Low Precision library developed by Google. Although this library is able to achieve overall better results due to the inclusion of auto-tuning approaches, it is possible that through inclusion of our approaches, the store, load, quantization and subsequently the auto-tuning approaches could achieve better overall results. On the other hand, the quantization process only needs to be performed once, and it is not necessary for quantization to be performed on an embedded device, so the benefit of including the work performed in this project is of minute importance in a general setting.

6. REFERENCES

- [1] Filippo Amato, Alberto Lpez, Eladia Mara Pea-Mndez, Petr Vahara, Ale Hampl, and Josef Havel, “Artificial neural networks in medical diagnosis,” *Journal of Applied Biomedicine*, vol. 11, no. 2, pp. 47–58, 2013.
- [2] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton, “ImageNet Classification with Deep Convolutional Neural Networks,” in *Advances in Neural Information Processing Systems 25*, F. Pereira, C. J. C. Burges, L. Bottou, and K. Q. Weinberger, Eds., pp. 1097–1105. Curran Associates, Inc., 2012.
- [3] Shixiang Gu, Ethan Holly, Timothy Lillicrap, and Sergey Levine, “Deep Reinforcement Learning for Robotic Manipulation with Asynchronous Off-Policy Updates,” *arXiv:1610.00633 [cs]*, Oct. 2016, arXiv: 1610.00633.
- [4] Yunchao Gong, Liu Liu, Ming Yang, and Lubomir Bourdev, “Compressing Deep Convolutional Networks using Vector Quantization,” *arXiv:1412.6115 [cs]*, Dec. 2014, arXiv: 1412.6115.
- [5] Wenlin Chen, James T. Wilson, Stephen Tyree, Kilian Q. Weinberger, and Yixin Chen, “Compressing Neural Networks with the Hashing Trick,” *arXiv:1504.04788 [cs]*, Apr. 2015, arXiv: 1504.04788.
- [6] Emily L Denton, Wojciech Zaremba, Joan Bruna, Yann LeCun, and Rob Fergus, “Exploiting Linear Structure Within Convolutional Networks for Efficient Evaluation,” in *Advances in Neural Information Processing Systems 27*, Z. Ghahramani, M. Welling, C. Cortes, N. D. Lawrence, and K. Q. Weinberger, Eds., pp. 1269–1277. Curran Associates, Inc., 2014.
- [7] Qinyao He, He Wen, Shuchang Zhou, Yuxin Wu, Cong Yao, Xinyu Zhou, and Yuheng Zou, “Effective Quantization Methods for Recurrent Neural Networks,” *arXiv:1611.10176 [cs]*, Nov. 2016, arXiv: 1611.10176.
- [8] Itay Hubara, Matthieu Courbariaux, Daniel Soudry, Ran El-Yaniv, and Yoshua Bengio, “Binarized Neural Networks,” in *Advances in Neural Information Processing Systems 29*, D. D. Lee, M. Sugiyama, U. V. Luxburg, I. Guyon, and R. Garnett, Eds., pp. 4107–4115. Curran Associates, Inc., 2016.
- [9] Vincent Vanhoucke, Andrew Senior, and Mark Z. Mao, “Improving the speed of neural networks on CPUs,” 2011.